

MACHINE LEARNING --- SAMPLE PAPER 1 --- IDEAL SOLUTION

Q1(b) --- Overfitting vs Underfitting

Overfitting occurs when the model learns noise in the training data, capturing patterns that do not generalize. **Underfitting** occurs when the model fails to capture the underlying structure of the data.

Aspect	Underfitting	Overfitting
Bias	High bias	Low bias
Variance	Low variance	High variance
Training error	High	Near zero
Test error	High	High (gap from training)
Example	Linear model on quadratic data	High-degree polynomial on small data
Cause	Model too simple	Model too complex

Regularization mitigates overfitting by adding a penalty term to the loss function: $L_{\text{total}} = L_{\text{original}} + \lambda \|w\|$. Ridge (L2) shrinks coefficients proportionally; Lasso (L1) can drive coefficients to zero, performing feature selection.

Q1(c) --- Regression Evaluation Metrics

MAE (Mean Absolute Error): $MAE = (1/n) \sum |y_i - \hat{y}_i|$

- Measures average absolute deviation
- Robust to outliers, linear penalty
- Same units as target variable

RMSE (Root Mean Squared Error): $RMSE = \sqrt{(1/n) \sum (y_i - \hat{y}_i)^2}$

- Measures standard deviation of residuals
- Quadratic penalty --- more sensitive to large errors
- Same units as target variable

R² (Coefficient of Determination): $R^2 = 1 - (SS_{\text{res}} / SS_{\text{tot}})$

- Proportion of variance explained by the model
- Ranges from $(-\infty, 1]$; 1 = perfect fit
- Can be artificially inflated by adding features (adjusted R² corrects this)

Metric	Range	Sensitivity to Outliers	Interpretation
MAE	$[0, \infty)$	Low	Average absolute error
RMSE	$[0, \infty)$	High	Standard deviation of errors
R ²	$(-\infty, 1]$	Depends	Goodness of fit

Q2(a) --- Lasso and Ridge Regression

Ridge Regression (L2): $L = \sum (y_i - \hat{y}_i)^2 + \lambda \sum w_j^2$

- Shrinks coefficients toward zero but never exactly to zero
- Handles multicollinearity well
- All features retained with reduced magnitude

Lasso Regression (L1): $L = \sum (y_i - \hat{y}_i)^2 + \lambda \sum |w_j|$

- Can drive coefficients exactly to zero (feature selection)
- Produces sparse models
- Preferred when many features are irrelevant

Effect of λ :

- $\lambda = 0$: standard OLS (may overfit)
- $\lambda \rightarrow \infty$: coefficients $\rightarrow 0$ (underfitting)
- Optimal λ found via cross-validation

```

flowchart LR
    subgraph Ridge["Ridge (L2)"]
        C1[Feature 1] -->|"β shrunk"| M1[Model]
        C2[Feature 2] -->|"β shrunk"| M1
        C3[Feature 3] -->|"β ≈ 0"| M1
    end
    subgraph Lasso["Lasso (L1)"]
        C4[Feature 1] -->|"β retained"| M2[Model]
        C5[Feature 2] -->|"β = 0"| X5[ ]
        C6[Feature 3] -->|"β retained"| M2
    end

```

Q2(b) --- Least Squares Regression (Numerical)

Given: $X = \begin{bmatrix} 2 & 4 & 6 & 8 & 10 \end{bmatrix}$ $Y = \begin{bmatrix} 55 & 60 & 68 & 75 & 82 \end{bmatrix}$

$n = 5$, $\sum X = 30$, $\sum Y = 340$, $\sum XY = 2140$, $\sum X^2 = 220$

$\bar{x} = 30/5 = 6$, $\bar{y} = 340/5 = 68$

$\beta_1 = (\sum XY - n \cdot \bar{x} \cdot \bar{y}) / (\sum X^2 - n \cdot \bar{x}^2)$ $\beta_1 = (2140 - 5 \times 6 \times 68) / (220 - 5 \times 36)$ $\beta_1 = (2140 - 2040) / (220 - 180) = 100/40 = 2.5$

$\beta_0 = \bar{y} - \beta_1 \bar{x} = 68 - 2.5 \times 6 = 68 - 15 = 53$

Regression line: $Y = 53 + 2.5X$

Prediction for $X = 12$: $Y = 53 + 2.5 \times 12 = 53 + 30 = 83$

Answer: $Y = 53 + 2.5X$, Estimated score for 12 hours = 83

Q2(c) --- Gradient Descent

Gradient Descent is an iterative optimization algorithm that minimizes the loss function $J(\theta)$ by moving in the direction of steepest descent: $\theta_j := \theta_j - \alpha \cdot (\partial/\partial \theta_j) J(\theta)$

Types of Gradient Descent:

Type	Data per Iteration	Noise	Convergence	Memory
Batch GD	Entire dataset	Low	Smooth, guaranteed	High
Stochastic GD	1 sample	High	Oscillating, escapes local minima	Low

Type	Data per Iteration	Noise	Convergence	Memory
Mini-batch GD	k samples (e.g., 32)	Moderate	Good balance	Moderate

Learning rate α controls step size:

- α too large: may diverge
- α too small: slow convergence
- Adaptive methods (Adam, RMSProp) adjust α per parameter

Unit IV --- Supervised Learning: Classification

Q3(a) --- K-Nearest Neighbors (KNN)

KNN is a non-parametric, lazy supervised learning algorithm that classifies a sample based on the majority class among its K nearest neighbors in the feature space.

Algorithm:

1. Choose K and a distance metric (Euclidean, Manhattan, Minkowski)
2. For each test sample, compute distances to all training samples
3. Select K nearest neighbors
4. Assign the majority class label

Example: Classify point (3, 4) with K=3. Training: Class A = {(1,2), (2,3)}, Class B = {(5,6), (7,8)}. Distances: to (1,2)= $\sqrt{8} \approx 2.8$, (2,3)= $\sqrt{2} \approx 1.4$, (5,6)= $\sqrt{8} \approx 2.8$, (7,8)= $\sqrt{41} \approx 6.4$. Nearest: (2,3)=A, (1,2)=A, (5,6)=B → **Class A**.

Effect of K: Small K = low bias, high variance (sensitive to noise). Large K = high bias, low variance (smoother boundaries). K is typically chosen as odd to avoid ties, tuned via cross-validation.

Q3(b) --- Ensemble Learning: Bagging and Boosting

Ensemble Learning combines multiple base learners to produce a model with better generalization than individual learners.

Bagging (Bootstrap Aggregating):

- Train multiple models on bootstrap samples (random sampling with replacement)
- Aggregate predictions: voting (classification), averaging (regression)
- Reduces **variance** without increasing bias
- **Random Forest** = Bagging + random feature subset selection at each split

Boosting:

- Train models sequentially, each focusing on misclassified samples from the previous model
- Weighted voting for final prediction
- Reduces **bias** and variance
- Examples: AdaBoost, Gradient Boosting, XGBoost

Aspect	Bagging	Boosting
Training	Parallel	Sequential
Bias	Unchanged	Reduced

Aspect	Bagging	Boosting
Variance	Reduced	Reduced
Overfitting	Less prone	Can overfit
Example	Random Forest	AdaBoost, XGBoost

Q3(c) --- Classifier Evaluation Metrics

Confusion Matrix:

	Predicted Positive	Predicted Negative
Actual Positive	TP (True Positive)	FN (False Negative)
Actual Negative	FP (False Positive)	TN (True Negative)

- **Accuracy** = $(TP + TN) / (TP + TN + FP + FN)$ --- overall correctness
- **Precision** = $TP / (TP + FP)$ --- how many selected items are relevant
- **Recall (Sensitivity)** = $TP / (TP + FN)$ --- how many relevant items are selected
- **F1-Score** = $2 \cdot (\text{Precision} \cdot \text{Recall}) / (\text{Precision} + \text{Recall})$ --- harmonic mean, balances precision and recall

When to use: F1-score is preferred for imbalanced datasets. Accuracy can be misleading when classes are skewed (e.g., 95% negative, classifying all as negative gives 95% accuracy but 0% recall for positive class).

Q4(a) --- Binary vs Multiclass Classification

Binary Classification: Two mutually exclusive classes (e.g., Spam/Not Spam). Output is a single probability thresholded at 0.5.

Multiclass Classification: Three or more classes (e.g., digit recognition 0-9).

Strategies for Multiclass:

Strategy	One-vs-All (OvA)	One-vs-One (OvO)
Classifiers	k binary classifiers	$k(k-1)/2$ binary classifiers
Training	Each class vs rest	Each pair of classes
Prediction	Highest confidence score	Majority voting
Scalability	Better for large k	Better for small k
Example	OvA: 10 classifiers for digits	OvO: 45 classifiers for digits

Q4(b) --- Support Vector Machine (SVM)

SVM is a supervised classification algorithm that finds the **maximum margin hyperplane** that best separates classes.

Key concepts:

- **Support Vectors:** Data points closest to the decision boundary that define the margin
- **Margin:** Distance between the hyperplane and the nearest support vectors

- **Hard Margin:** Perfectly separable data
- **Soft Margin:** Allows misclassifications (controlled by C parameter)

Kernel Trick: Maps data to higher-dimensional space without explicit computation:

- Linear: $K(x,y) = x \cdot y$
- Polynomial: $K(x,y) = (x \cdot y + c)^d$
- RBF (Gaussian): $K(x,y) = \exp(-\gamma \|x-y\|^2)$

Flowchart TD

```
subgraph Input["Input Space"]
```

```
  P1["Data Points"]
```

```
end
```

```
P1 -->| "Kernel Function  $\phi(x)$ " | FS["Feature Space<br/>(Higher Dimension)"]
```

```
FS -->| "Maximum Margin<br/>Hyperplane" | Decision["Class +1 or -1"]
```

Q4(c) --- Handling Imbalanced Data

Imbalanced data occurs when classes have significantly different sample sizes. Common techniques:

1. **Random Oversampling:** Duplicate samples from minority class
2. **Random Undersampling:** Remove samples from majority class
3. **SMOTE (Synthetic Minority Over-sampling Technique):** Create synthetic samples by interpolating between minority class neighbors
4. **Cost-sensitive Learning:** Assign higher misclassification cost to minority class
5. **Ensemble Methods:** Balanced Random Forest, EasyEnsemble

Method	Advantage	Disadvantage
Oversampling	No information loss	Risk of overfitting
Undersampling	Reduces training time	Loses potentially useful data
SMOTE	Creates realistic synthetic samples	May create noisy samples
Cost-sensitive	No data modification	Requires domain-specific cost

Unit V --- Unsupervised Learning

Q5(a) --- K-Means Clustering

K-Means is a partition-based clustering algorithm that groups data into K clusters by minimizing within-cluster variance.

Steps:

1. Initialize K centroids
2. Assign each point to nearest centroid
3. Recompute centroids as mean of assigned points
4. Repeat steps 2-3 until convergence

Given points: A(1,2), B(2,1), C(4,5), D(5,4), E(7,8) **Initial centroids:** C1 = (1,1), C2 = (5,5)

Iteration 1:

- Distances to C1: $A(\sqrt{1}=1)$, $B(\sqrt{1}=1)$, $C(\sqrt{25}=5)$, $D(\sqrt{25}=5)$, $E(\sqrt{85}=9.2)$
- Distances to C2: $A(\sqrt{25}=5)$, $B(\sqrt{25}=5)$, $C(\sqrt{1}=1)$, $D(\sqrt{1}=1)$, $E(\sqrt{13}=3.6)$
- Cluster 1: {A, B}, Cluster 2: {C, D, E}
- New C1 = $((1+2)/2, (2+1)/2) = (1.5, 1.5)$
- New C2 = $((4+5+7)/3, (5+4+8)/3) = (5.33, 5.67)$

Iteration 2:

- Distances to C1(1.5,1.5): $A(\sqrt{0.5}=0.7)$, $B(\sqrt{0.5}=0.7)$, $C(\sqrt{26}=5.1)$, $D(\sqrt{24.5}=4.95)$, $E(\sqrt{67.3}=8.2)$
- Distances to C2(5.33,5.67): $A(\sqrt{32.2}=5.7)$, $B(\sqrt{33.2}=5.8)$, $C(\sqrt{1.1}=1.05)$, $D(\sqrt{2.8}=1.67)$, $E(\sqrt{8.2}=2.86)$
- Cluster 1: {A, B}, Cluster 2: {C, D, E} --- **No change, converged**

Answer: Cluster 1 = {A, B} centroid (1.5, 1.5), Cluster 2 = {C, D, E} centroid (5.33, 5.67)

Q5(b) --- Outlier Analysis

Outlier Analysis is the process of identifying data points that deviate significantly from the majority of data.

Importance: Outliers can indicate errors (data entry, sensor malfunction), fraud (credit card transactions), or novel discoveries (medical anomalies).

Detection Methods:

- **Isolation Forest:** Randomly partitions data; outliers require fewer partitions to isolate
- **Local Outlier Factor (LOF):** Measures local density deviation --- outliers have lower density than neighbors
- **Z-Score:** Points beyond 3 standard deviations from mean
- **IQR:** Points below $Q1 - 1.5 \times IQR$ or above $Q3 + 1.5 \times IQR$

Method	Advantage	Disadvantage
Z-Score	Simple, interpretable	Assumes normality
IQR	Distribution-free	Ignores local patterns
LOF	Detects local outliers	Sensitive to parameter k
Isolation Forest	Efficient for high dimensions	Random (non-deterministic)

Q5(c) --- Elbow Method

Elbow Method is a heuristic for determining the optimal number of clusters K in K-Means.

Procedure:

1. Compute K-Means for $K = 1$ to K_{max}
2. Plot **Within-Cluster Sum of Squares (WCSS)** against K
3. Choose K where WCSS starts to decrease linearly (the "elbow point")

Silhouette Score measures how similar a point is to its own cluster vs neighboring clusters, ranging from -1 to 1. Higher values indicate better-defined clusters. It complements the elbow method by providing a quantitative validation of cluster quality.

Q6(a) --- Hierarchical Clustering

Agglomerative Hierarchical Clustering builds a hierarchy by iteratively merging the closest pair of clusters.

Linkage criteria:

- **Single-linkage:** Distance between closest members of clusters --- produces chain-like clusters
- **Complete-linkage:** Distance between farthest members --- produces compact clusters
- **Average-linkage:** Average of all pairwise distances --- balances the two

```
Flowchart TD
    subgraph Dendrogram["Dendrogram"]
        A[Point A] --> M1[Merge]
        B[Point B] --> M1
        M1 --> M2[Merge]
        C[Point C] --> M3[Merge]
        D[Point D] --> M3
        M3 --> M2
    end
    end
```

Q6(b) --- DBSCAN

DBSCAN (Density-Based Spatial Clustering of Applications with Noise) groups points that are closely packed together, marking points in low-density regions as noise.

Parameters: ϵ (neighborhood radius), minPts (minimum points to form a dense region)

Properties:

- Can find arbitrarily shaped clusters
- Automatically identifies noise points
- No need to specify number of clusters
- Struggles with varying density clusters

Aspect	K-Means	DBSCAN
Shape	Spherical only	Arbitrary shapes
Noise	Forces all points to clusters	Explicit noise handling
K	Must specify	Automatic
Parameters	K	ϵ , minPts

Unit VI --- Introduction to Neural Networks

Q7(a) --- Activation Functions

Activation functions introduce non-linearity into neural networks, enabling them to learn complex patterns.

Sigmoid: $f(x) = 1/(1+e^{-x})$, range (0,1) --- used in binary classification output **Tanh:** $f(x) = 2\sigma(2x)-1 = (e^x-e^{-x})/(e^x+e^{-x})$, range (-1,1) --- zero-centered, better gradient flow **ReLU:** $f(x) = \max(0,x)$, range $[0,\infty)$ --- most common in hidden layers, avoids vanishing gradient **Softmax:** $f(x_i) = e^{x_i}/\sum_j e^{x_j}$ --- multiclass probability output

Function	Range	Gradient	Use Case
Sigmoid	(0,1)	Vanishes for large x	Binary output
Tanh	(-1,1)	Vanishes but zero-centered	Hidden layers
ReLU	$[0, \infty)$	0 or 1 (no vanish)	Hidden layers (default)
Softmax	(0,1)	Full Jacobian	Multiclass output

Q7(b) --- Recurrent Neural Networks and LSTM

RNN is a neural network architecture designed for sequential data, maintaining a hidden state that captures information from previous time steps: $h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t + b_h)$

Example: Language modeling --- given "the cat sat on the ____", RNN predicts "mat" using context from previous words.

Vanishing Gradient Problem: During backpropagation through time (BPTT), gradients can exponentially diminish, preventing the network from learning long-range dependencies.

LSTM (Long Short-Term Memory) addresses this with:

- **Forget gate:** Decides what to discard from cell state
- **Input gate:** Decides what new information to store
- **Output gate:** Decides what to output based on cell state
- **Cell state:** Maintains information over long sequences

```

flowchart LR
    subgraph LSTM["LSTM Cell"]
        F[Forget Gate] --> C[Cell State]
        I[Input Gate] --> C
        C --> O[Output Gate]
    end
    end
    X_prev["h(t-1)"] --> LSTM
    X_t["x(t)"] --> LSTM
    LSTM --> H_t["h(t)"]
    LSTM --> C_t["c(t)"]

```

Q7(c) --- CNN Architecture

```

flowchart LR
    Input["Input Image<br/>32x32x3"] --> Conv1["Convolutional Layer<br/>Filters: 6x5x5<br/>$\to $ 28x28x6"]
    Conv1 --> ReLU1["ReLU Activation"]
    ReLU1 --> Pool1["Pooling Layer<br/>Max Pool 2x2<br/>$\to $ 14x14x6"]
    Pool1 --> Conv2["Convolutional Layer<br/>Filters: 16x5x5<br/>$\to $ 10x10x16"]
    Conv2 --> ReLU2["ReLU Activation"]
    ReLU2 --> Pool2["Pooling Layer<br/>2x2 $\to $ 5x5x16"]
    Pool2 --> FC1["Fully Connected<br/>120 neurons"]
    FC1 --> FC2["Fully Connected<br/>84 neurons"]
    FC2 --> Output["Output Layer<br/>10 classes (Softmax)"]

```

Convolutional Layer: Applies learnable filters (kernels) to detect features (edges, textures). Stride and padding control output dimensions.

Pooling Layer: Reduces spatial dimensions, providing translation invariance. Max pooling retains strongest activation; average pooling smooths features.

Fully Connected Layer: Combines extracted features for final classification.

Q8(a) --- Backpropagation vs Feedforward

Feedforward Network: Information flows in one direction --- input → hidden → output. No feedback connections. Used for static pattern mapping.

Backpropagation Network: Same architecture but trained using backpropagation --- an algorithm that computes gradients of the loss with respect to weights using the **chain rule**.

Aspect	Feedforward	Backpropagation
Direction	Forward only	Forward + backward
Weight update	Not applicable	Gradient descent
Learning	Requires external algorithm	Integrated training
Purpose	Inference	Training

Backpropagation algorithm:

1. Forward pass: compute predictions and loss
 2. Backward pass: compute $\partial L / \partial w$ for each layer
 3. Update: $w = w - \eta \cdot \partial L / \partial w$
-

Q8(b) --- Padding in CNNs

Padding adds extra pixels around the input boundary to control the output spatial dimensions.

Valid Padding: No padding. Output size = $n - f + 1$ (where n =input size, f =filter size). Output is smaller than input.

Same Padding: Pad such that output size equals input size. $p = (f-1)/2$. Zero-padding is most common.

Effect: Without padding, each convolution shrinks the spatial dimensions and loses edge information. Same padding preserves size and edge features.

Q8(c) --- Convolutional and Pooling Layers

i) Convolutional Layer: Applies a set of learnable filters (kernels) that slide across the input. Each filter detects specific features --- early layers detect edges, later layers detect complex patterns. Parameters: filter size, stride, padding, number of filters.

ii) Pooling Layer: Down-samples feature maps:

- **Max Pooling:** Outputs maximum value in each window --- preserves strongest features
- **Average Pooling:** Outputs average value --- preserves overall information
- **Global Pooling:** Reduces entire feature map to a single value

Pooling reduces computational load, controls overfitting, and provides translation invariance.

EXAMINER COMMENTARY Why this scores

full marks: Each answer follows the definition → explanation → example/closing structure.

Numerical parts show all steps with boxed answers. Comparisons use tables with minimum 3 bases. Diagrams are embedded where structural. **Common Deductions:**

- Regression derivation without showing gradient steps or partial derivatives
 - K-Means without showing distance calculations for each iteration
 - Confusion matrix without explanation of TP/TN/FP/FN
 - Activation functions listed without mathematical expressions
 - CNN architecture drawn without labeling all layers
- Time Budget:**
- Q1/Q2 (18 marks): 42 min
 - Q3/Q4 (17 marks): 40 min
 - Q5/Q6 (18 marks): 42 min
 - Q7/Q8 (17 marks): 38 min
-
-