

Unit III --- Convolution Neural Network

Q1(a) --- CNN Architecture

CNN is a deep neural network architecture designed for processing grid-like data (images), using convolution operations to automatically learn spatial hierarchies of features.

flowchart LR

```
Input["Input Image<br/>32×32×3"] --> Conv1["Conv Layer 1<br/>6 filters × 5×5<br/>Stride 1, Same pad<br/>→ 32×32×6"]
Conv1 --> Act1["ReLU"]
Act1 --> Pool1["Max Pool<br/>2×2, Stride 2<br/>→ 16×16×6"]
Pool1 --> Conv2["Conv Layer 2<br/>16 filters × 5×5<br/>→ 16×16×16"]
Conv2 --> Act2["ReLU"]
Act2 --> Pool2["Max Pool<br/>2×2, Stride 2<br/>→ 8×8×16"]
Pool2 --> Flat["Flatten<br/>→ 1024"]
Flat --> FC1["FC Layer<br/>120 neurons"]
FC1 --> FC2["FC Layer<br/>84 neurons"]
FC2 --> Out["Output<br/>10 classes (Softmax)"]
```

Role of each layer:

1. **Input Layer:** Raw pixel values (height × width × channels)
2. **Convolutional Layer:** Learns feature maps by sliding filters across input
3. **Activation Layer (ReLU):** Introduces non-linearity
4. **Pooling Layer:** Reduces spatial dimensions, provides translation invariance
5. **Fully Connected Layer:** Combines high-level features for classification
6. **Output Layer:** Softmax produces class probabilities

Q1(b) --- Working of Convolution Layer

Convolution Layer applies learnable filters (kernels) that slide across the input to produce feature maps.

Operation: Each filter computes the dot product between its weights and a local region of the input. Multiple filters produce multiple feature maps.

$$\text{Output dimension formula: } O = \lfloor (n - f + 2p) / s \rfloor + 1$$

Where n = input size, f = filter size, p = padding, s = stride

Feature extraction: Early layers detect low-level features (edges, corners). Deeper layers combine these into high-level features (shapes, objects, faces).

Stride controls how much the filter shifts each step. Larger stride = smaller output.

Padding adds zeros around the input to preserve spatial dimensions.

- Valid padding ($p=0$): output shrinks
- Same padding ($p=(f-1)/2$): output same size as input

Q1(c) --- Pooling Layers

Pooling is a down-sampling operation that reduces spatial dimensions and provides translation invariance.

Types:

- **Max Pooling:** Takes maximum value in each window. Preserves strongest activations. Most common.
- **Average Pooling:** Takes average value. Smoother, preserves background information.
- **Global Pooling:** Reduces entire feature map to a single value. Used before final FC layers.

Type	Operation	Preservation	Use Case
Max Pool	Maximum	Edge/texture features	Hidden layers
Average Pool	Mean	Overall signal	Transition layers
Global Pool	Mean over entire map	Channel importance	Before output

Max Pooling is preferred in most architectures because it selects the most important features and provides better translation invariance.

Q2(a) --- Features of Pooling Layer

Pooling Layer provides:

- Dimensionality reduction:** Reduces spatial size, lowering computational cost
- Translation invariance:** Small shifts in input produce same pooled output
- Overfitting prevention:** Fewer parameters reduce model capacity
- Receptive field increase:** Subsequent layers see larger portions of input
- Feature selection:** Max pooling retains dominant features

Q2(b) --- Local Response Normalization (LRN)

LRN is a normalization technique that creates competition among neuron outputs computed by different filters at the same spatial location.

$$\text{Formula: } b_i^{xy} = a_i^{xy} / (k + \alpha \sum_{j=\max(0, i-n/2)}^{\min(N-1-i+n/2)} (a_j^{xy})^2)^\beta$$

Where a_i^{xy} is the output of filter i at position (x,y) , and N is the total number of filters.

Effect: LRN encourages lateral inhibition --- strongly activated neurons suppress nearby neurons, promoting feature diversity.

Aspect	LRN	Batch Normalization
Scope	Across channels	Across batch
Normalization	Per spatial location	Per feature
Location	After activation	Before or after activation
Modern usage	Rare (AlexNet era)	Standard in modern CNNs

Batch Normalization has largely replaced LRN as it enables higher learning rates, reduces internal covariate shift, and provides regularization.

Q2(c) --- ReLU Layer

ReLU (Rectified Linear Unit): $f(x) = \max(0, x)$

Advantages over Sigmoid/Tanh:

- Computational efficiency:** Simple threshold operation vs exponential
- No vanishing gradient:** Gradient = 1 for positive inputs, 0 for negative
- Sparse activation:** Only positive neurons fire
- Faster convergence:** Empirically 6× faster than Tanh

Dying ReLU problem: When many inputs are negative, ReLU neurons can permanently die (output = 0, gradient = 0) for all inputs, never recovering.

Solution --- Leaky ReLU: $f(x) = \max(\alpha x, x)$ where α is a small constant (e.g., 0.01). Provides a small gradient for negative inputs, preventing dead neurons.

Unit IV --- Recurrent Neural Network

Q3(a) --- Recursive vs Recurrent Neural Networks


Recursive Neural Network operates on hierarchical structures (trees) by recursively applying the same weight matrix to child nodes. Used for natural language parsing and compositional semantics.

RNN operates on sequential data (chains) by applying the same transition function at each time step.

Aspect	Recursive NN	Recurrent NN
Structure	Tree	Chain
Data type	Hierarchical (parse trees)	Sequential (time series)
Weight sharing	Across nodes in tree	Across time steps
Application	Sentiment analysis of phrases	Language modeling
Backprop	Through structure (BPTS)	Through time (BPTT)

Q3(b) --- LSTM Architecture

LSTM (Long Short-Term Memory) is a gated RNN architecture designed to capture long-range dependencies by controlling information flow through three gates.



Forget gate: Decides which information to discard from cell state
Input gate: Decides which new information to store in cell state
Output gate: Decides what to output based on cell state

Q3(c) --- Working of RNN

RNN processes sequential data by maintaining a hidden state that captures information from previous time steps.

$$\text{Hidden state update: } h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t + b_h)$$

The hidden state h_t acts as memory, carrying information across time steps. It is updated at each step based on the previous hidden state and current input.

Variable-length sequences: RNNs handle variable-length inputs by processing one token at a time, updating the hidden state. The final hidden state (or all hidden states) can be used for downstream tasks. Padding or bucketing handles batch processing of different-length sequences.

Q4(a) --- CNN vs RNN

Aspect	CNN	RNN
Data	Grid (images, 2D/3D)	Sequence (text, time series)
Connection	Feedforward	Recurrent (feedback loops)
Weight sharing	Across spatial locations	Across time steps
Memory	No internal state	Hidden state as memory
Parallelization	High (independent patches)	Low (sequential dependence)
Use case	Image classification, object detection	Language modeling, translation
Gradient flow	Stable	Vanishing/exploding (BPTT)

Q4(b) --- Long-Term Dependencies and Gradient Problems

Long-term dependencies: RNNs struggle to learn patterns separated by many time steps because gradient signals decay or explode through repeated multiplication.

Vanishing gradients: During backpropagation through time (BPTT), gradients are multiplied by W_{hh} at each step. If W_{hh} 's eigenvalues < 1 , gradients $\rightarrow 0$ exponentially, and early time steps receive negligible gradient updates.

Exploding gradients: If eigenvalues > 1 , gradients $\rightarrow \infty$. Addressed by gradient clipping.

Gated architectures (LSTM/GRU) alleviate vanishing gradients through:

- **Cell state:** Direct linear connection across time (not squashed by activation)
- **Gates:** Learn to control information flow (forget irrelevant information, retain relevant)
- **Additive gradients:** Error signals flow through the cell state without repeated non-linear transformations

Q4(c) --- Encoder-Decoder with Attention

Encoder-Decoder (Seq2Seq): An architecture for mapping variable-length input sequences to variable-length output sequences.

- **Encoder RNN:** Reads input sequence and produces a context vector (final hidden state)
- **Decoder RNN:** Generates output sequence conditioned on the context vector and previous outputs

Machine translation: English "I am a student" $\rightarrow$ Encoder $\rightarrow$ context vector $\rightarrow$ Decoder $\rightarrow$ "Je suis étudiant"

Attention mechanism: Instead of compressing the entire input into a single context vector, attention allows the decoder to focus on relevant parts of the input at each step. Attention weights compute alignment between decoder state and all encoder hidden states.

```
flowchart LR
    subgraph Encoder["Encoder"]
        E1[EOS] --> E2[ENG]
        E2 --> E3[ENC]
        E3 --> E4[EOS]
    end
    E4 --> C[Context Vector]
    subgraph Decoder["Decoder with Attention"]
        C --> D1[BOS]
        D1 --> D2[FRE]
        D2 --> D3[DEC]
        D3 --> D4[EOS]
    end
    E2 -. -> |"Attention α₁"| D2
    E3 -. -> |"Attention α₂"| D2
```

Unit V --- Deep Generative Models

Q5(a) --- Deep Generative Models

Deep Generative Models learn the joint probability distribution $P(X, Y)$ of the data, enabling them to generate new samples from the learned distribution.

Difference from discriminative models: Discriminative models learn $P(Y|X)$ --- the decision boundary. Generative models learn the full data distribution $P(X)$, enabling sampling of new data.

Applications: Image generation (DALL-E, Stable Diffusion), data augmentation (generating synthetic training data), anomaly detection, super-resolution, text-to-image synthesis.

Q5(b) --- Boltzmann Machine

Boltzmann Machine is a stochastic recurrent neural network where neurons are connected symmetrically, learning internal representations from binary data.

$$\text{Energy function: } E(\mathbf{v}, \mathbf{h}) = -\sum_i v_i b_i - \sum_j h_j c_j - \sum_i \sum_j v_i w_{ij} h_j$$

Where $\mathbf{v}$ = visible units, $\mathbf{h}$ = hidden units, w = connection weights, b , c = biases.

Training: Contrastive Divergence minimizes the energy of observed data while maximizing the energy of unobserved configurations.

RBM (Restricted Boltzmann Machine) removes connections between visible-visible and hidden-hidden, forming a bipartite graph. This enables efficient Gibbs sampling and tractable inference. RBMs can be stacked to form Deep Belief Networks.

Aspect	Boltzmann Machine	Restricted BM
Connections	Visible-visible, hidden-hidden, visible-hidden	Only visible-hidden
Inference	Intractable for large networks	Tractable (conditional independence)
Training	Slow (MCMC)	Efficient (CD-k)
Use	Theoretical foundation	Building block for DBNs

Q5(c) --- Generative Adversarial Network (GAN)

GAN consists of two networks trained adversarially: a Generator G that creates fake data, and a Discriminator D that distinguishes real from fake.

```
flowchart LR
    Z["Z[Noise Vector z]"] --> G["G[Generator G]"]
    G --> Fake["Fake[Fake Samples]"]
    Real["Real[Real Data]"] --> D["D[Discriminator D]"]
    Fake --> D
    D --> Output["Output[Real / Fake]"]
    Output --> Loss["|\"Loss\"| G"]
```

Minimax objective: $\min_G \max_D V(D, G) = E_{x \sim p_{data}}[\log D(x)] + E_{z \sim p_z}[\log(1 - D(G(z)))]$

Training process: Alternate between updating D (distinguish real from fake) and G (fool D).

Convergence occurs when G generates data indistinguishable from real and D outputs 1/2.

Example --- MNIST generation: G takes random noise (100-dim) $\rightarrow$ deconvolution layers $\rightarrow$ 28×28 image.

D takes 28×28 image $\rightarrow$ convolution layers $\rightarrow$ binary classification.

Q6(a) --- Deep Belief Networks (DBN)

DBN is a generative model composed of stacked RBMs, where each RBM's hidden layer serves as the visible layer for the next RBM.

Greedy layer-wise pretraining:

1. Train RBM_1 (visible = input, hidden = h_1) using Contrastive Divergence
2. Freeze RBM_1 weights, train RBM_2 (visible = h_1 , hidden = h_2)
3. Repeat for each layer

4. Fine-tune all layers using supervised backpropagation

```
flowchart TD
    subgraph DBN["Deep Belief Network"]
        V[Visible Layer] <--> H1[Hidden Layer 1<br/>RBM1]
        H1 <--> H2[Hidden Layer 2<br/>RBM2]
        H2 <--> H3[Hidden Layer 3<br/>RBM3]
        H3 --> Output[Output Layer]
    end
end
```

Benefits: Pretraining initializes weights near a good solution, preventing poor local minima and enabling deeper networks.

Q6(b) --- GAN Components

Generator $G(z)$: Maps latent noise z (random vector) to data space. Uses transposed convolutions. Learns to produce realistic samples. Goal: maximize $D(G(z))$ such that D classifies it as real.

Discriminator $D(x)$: Binary classifier distinguishing real from generated data. Uses convolutional layers. Outputs probability of realness. Goal: maximize $D(x)$ for real x , minimize $D(G(z))$.

Training loop: For each iteration: (1) Sample real batch + noise batch. (2) Update D : $\nabla_{\theta} D[\log D(x) + \log(1 - D(G(z)))]$. (3) Update G : $\nabla_{\theta} G[\log D(G(z))]$. Repeat until Nash equilibrium.

Q6(c) --- Types of GAN

DCGAN (Deep Convolutional GAN): Uses convolutional layers (instead of MLPs) for both G and D . Introduces batch normalization, strided convolutions, and no fully connected layers. Significantly improves training stability.

Conditional GAN (cGAN): Both G and D receive additional conditioning information (class labels, text). Enables controlled generation: generate images of a specific class.

CycleGAN: Enables image-to-image translation without paired examples (e.g., zebra $\leftrightarrow$ horse). Uses cycle consistency loss: translating zebra $\rightarrow$ horse $\rightarrow$ zebrafish should reconstruct the original.

GAN Type	Input	Key Feature	Application
DCGAN	Noise + Convolutions	Stable convolutional architecture	Image generation
cGAN	Noise + Condition	Controlled generation	Text-to-image
CycleGAN	Image A	Cycle consistency loss	Style transfer

Unit VI --- Reinforcement Learning

Q7(a) --- Markov Decision Process (MDP)

MDP is a mathematical framework for modeling sequential decision-making problems.

Components: (S, A, P, R, γ)

- S : Set of states (environment configurations)
- A : Set of actions (agent choices)
- $P(s'|s,a)$: Transition probability --- probability of reaching s' from s after action a
- $R(s,a,s')$: Reward function --- immediate reward received
- $\gamma \in [0,1]$: Discount factor (weights near vs far rewards)

Policy $\pi(a|s)$: Maps states to action probabilities.

Value Iteration: Iteratively updates state values: $V_{k+1}(s) = \max_a \sum_{s'} P(s'|s,a)[R(s,a,s') + \gamma V_k(s')]$, converging to optimal V^* .

Policy Iteration: Alternates between policy evaluation (computing V for current π) and policy improvement (updating π greedily with respect to V).

Q7(b) --- Deep Reinforcement Learning

Deep RL combines deep neural networks with RL to handle high-dimensional state spaces (e.g., raw pixels from games) where traditional tabular RL is infeasible.

Limitations addressed: Traditional RL requires discrete state spaces and hand-crafted features.

Deep RL learns feature representations automatically from raw input.

DQN (Deep Q-Network): Uses a CNN to approximate the Q-function $Q(s,a; \theta)$, taking raw game frames as input and outputting Q-values for each action.

Q7(c) --- Challenges of RL

1. **Exploration-exploitation tradeoff:** Agent must balance exploring unknown actions (gathering information) vs exploiting known high-reward actions. ϵ -greedy, UCB, and Thompson sampling address this.
2. **Delayed rewards:** Actions may affect rewards many steps later (e.g., winning a chess game). Credit assignment is difficult. Discounted returns help but don't fully solve it.
3. **Sample efficiency:** Deep RL requires millions of interactions with the environment. Model-based RL and offline RL aim to improve sample efficiency.
4. **Reward design:** Poorly designed reward functions can lead to unintended behaviors (reward hacking).

Q8(a) --- Deep Q-Learning

Deep Q-Learning approximates the optimal action-value function $Q^*(s,a)$ using a neural network.

Q-learning update rule: $Q(s,a) \leftarrow Q(s,a) + \alpha[r + \gamma \max_{a'} Q(s',a') - Q(s,a)]$

DQN improvements:

1. **Experience replay:** Stores transitions (s,a,r,s') in a replay buffer. Samples mini-batches randomly during training. Breaks correlation between consecutive samples, improving stability.
2. **Target network:** A separate network $\hat{Q}$ (with fixed parameters θ^-) computes target Q-values. θ^- is periodically updated to match the online network θ . Reduces oscillations during training.

flowchart TD

```
Env[Environment] -->|" (s,a,r,s') "| Buffer[Replay Buffer]
Buffer -->|"Random Batch"| Online[Online Network<br/>Q(s,a;θ)]
Online -->|"Loss L = (r + γQ̂(s',a') - Q(s,a))²"| Update[Update θ]
Update -->|"Every C steps"| Target[Target Network<br/>Q̂(s,a;θ⁻)]
Target --> LossCalc[Compute Target]
```

Q8(b) --- RL for Tic-Tac-Toe

Tic-Tac-Toe modeled as MDP:

- **States:** Board configurations (3×3 grid, each cell = empty/X/O)
- **Actions:** Place a marker in an empty cell
- **Transitions:** Deterministic (opponent's move is part of environment dynamics)
- **Rewards:** +1 for win, -1 for loss, 0 for draw, -0.1 for each move (encourage shorter wins)
- **Discount factor:** $\gamma = 1$ (finite horizon)

Learning process: Agent plays against itself (self-play) or a fixed opponent. Q-values are updated using the Q-learning rule after each move. After training (~10000 games), the agent converges to optimal play (never loses).

Q8(c) --- Dynamic Programming for RL

Dynamic Programming (DP) methods solve MDPs when the model (transition probabilities and rewards) is fully known.

Policy Evaluation: Compute $V_{\pi}(s)$ for a given policy π using *Bellman expectation backup*: $V_{k+1}(s) = \sum_a \pi(a|s) \sum_{s'} P(s'|s,a) [R(s,a,s') + \gamma V_k(s')]$

Policy Improvement: Generate a better policy π' from $V\pi$: $\pi'(s) = \operatorname{argmax}_a \sum \{s'\}$
 $P(s'|s,a)[R(s,a,s') + \gamma V_{\pi}(s')]$

Policy Iteration: Eval $\rightarrow$ Improve $\rightarrow$ Eval $\rightarrow$ Improve ... until convergence to optimal π^* .

EXAMINER

COMMENTARY Why this scores

full marks: Architecture diagrams are comprehensive with all layers labeled. LSTM cell diagram shows internal gate structure. Code-equivalent logic is explained stepwise. Mathematical formulations (energy functions, Q-learning updates) are included. Comparison tables have minimum 3 bases. **Common Deductions:**

- CNN architecture drawn without fully connected layers
 - LSTM explained textually without gate diagram
 - GAN training described without the minimax objective
 - DBN construction without explaining greedy pretraining
 - RL algorithms described without Bellman equation reference
- Time Budget:**
- Q1/Q2 (18 marks): 42 min
 - Q3/Q4 (17 marks): 40 min
 - Q5/Q6 (18 marks): 42 min
 - Q7/Q8 (17 marks): 38 min
-
-