

DATABASE MANAGEMENT SYSTEMS --- SAMPLE SOLUTION

Paper Code: -35 | Total Marks: 70 | Time: 2½ Hours

[6262]

Q1) a) Functional Dependency and Armstrong's Axioms

[9]

Functional Dependency (FD) is a constraint that describes the relationship between attributes in a relation. It is denoted as $X \rightarrow Y$, where X and Y are sets of attributes. This means that if two tuples have the same value for X , they must have the same value for Y .

Armstrong's Axioms are sound and complete inference rules used to derive all functional dependencies from a given set:

1. **Reflexivity:** If $Y \subseteq X$, then $X \rightarrow Y$
 - Example: $\{A, B\} \rightarrow A$
2. **Augmentation:** If $X \rightarrow Y$, then $XZ \rightarrow YZ$ for any Z
 - Example: If $A \rightarrow B$, then $AC \rightarrow BC$
3. **Transitivity:** If $X \rightarrow Y$ and $Y \rightarrow Z$, then $X \rightarrow Z$
 - Example: If $A \rightarrow B$ and $B \rightarrow C$, then $A \rightarrow C$

Additional derived rules:

- **Union:** If $X \rightarrow Y$ and $X \rightarrow Z$, then $X \rightarrow YZ$
- **Decomposition:** If $X \rightarrow YZ$, then $X \rightarrow Y$ and $X \rightarrow Z$
- **Pseudo-transitivity:** If $X \rightarrow Y$ and $WY \rightarrow Z$, then $WX \rightarrow Z$

Q1) b) Normalization Problem

[9]

Given: $R(A, B, C, D, E)$ with FDs: $A \rightarrow BC$, $CD \rightarrow E$, $B \rightarrow D$, $E \rightarrow A$

i) Finding Candidate Keys:

- **Closure of A:** $A^+ = A \rightarrow BC \rightarrow BCD$ (via $B \rightarrow D$) $\rightarrow BCDE$ (via $CD \rightarrow E$) $\rightarrow ABCDE$ ✓
- **Closure of E:** $E^+ = E \rightarrow A$ (via $E \rightarrow A$) $\rightarrow ABCDE$ ✓
- **Closure of CD:** $CD^+ = CD \rightarrow E \rightarrow A \rightarrow BC \rightarrow D \dots \rightarrow ABCDE$ ✓

Candidate Keys: A, E, CD

ii) Highest Normal Form:

Check for BCNF --- every FD's LHS must be a superkey:

- $A \rightarrow BC$: A is a candidate key ✓
- $CD \rightarrow E$: CD is a candidate key ✓
- $B \rightarrow D$: B is **not** a superkey ✗ $\rightarrow R$ is **not in BCNF**
- $E \rightarrow A$: E is a candidate key ✓

Check for 3NF --- for each $X \rightarrow Y$, either X is a superkey or Y is prime:

- $B \rightarrow D$: B is not a superkey but D is not prime ✗ $\rightarrow R$ is **not in 3NF**

Highest Normal Form: 2NF (no partial dependencies exist since no composite key)

iii) Decomposition into 3NF:

Step 1: Canonical cover $F_c = \{A \rightarrow B, A \rightarrow C, CD \rightarrow E, B \rightarrow D, E \rightarrow A\}$ Step 2: $R_1(A, B, C)$, $R_2(C, D, E)$, $R_3(B, D)$, $R_4(E, A)$ Step 3: $R_5(A, C)$ not needed as A already in R_1 Step 4: Check lossless --- yes, via shared attributes

Final 3NF Decomposition: $R_1(A, B, C)$, $R_2(B, D)$, $R_3(C, D, E)$, $R_4(E, A)$

[ANSWER BOX]

Final Decomposition: $R_1(A, B, C)$, $R_2(B, D)$, $R_3(C, D, E)$, $R_4(E, A)$

All in 3NF, lossless and dependency-preserving

Q2) a) Normalization Forms

[8]

Normalization is the process of organizing data in a database to reduce redundancy and eliminate anomalies (insertion, update, deletion).

1. **1NF: Each cell contains a single atomic value; no repeating groups.**
 - *Bad*: Student(Id, Name, Subjects) --- Subjects = {DBMS, TOC}
 - *Good*: Student(Id, Name, Subject) --- one row per subject
2. **2NF: In 1NF + no partial dependency (non-key attr must depend on full primary key).**
 - *Bad*: Enrollment(StudentId, CourseId, StudentName, Grade) --- StudentName depends only on StudentId
 - *Decompose*: Student(StudentId, Name), Enrollment(StudentId, CourseId, Grade)
3. **3NF: In 2NF + no transitive dependency (non-key attr must not depend on another non-key attr).**
 - *Bad*: Employee(EmpId, DeptId, DeptLocation) --- DeptLocation depends on DeptId
 - *Decompose*: Employee(EmpId, DeptId), Department(DeptId, DeptLocation)
4. **BCNF: In 3NF + every determinant must be a candidate key.**
 - *Bad*: StudentCourse(StudentId, CourseId, ProfessorName) where ProfessorName $\rightarrow$ CourseId
 - *Decompose*: StudentCourse(StudentId, ProfessorName), CourseProfessor(ProfessorName, CourseId)

Q2) b) Lossless Decomposition Check

[9]

Given: $R(A, B, C, D, E, F)$ with FDs: $A \rightarrow B$, $C \rightarrow D$, $D \rightarrow E$, $B \rightarrow F$

Decomposition: $R_1(A, B, F)$, $R_2(C, D, E)$, $R_3(A, C)$

Chase Algorithm:

	A	B	C	D	E	F
R1	a	b	c1	d1	e1	f
R2	a2	b2	c	d	e	f2
R3	a	b3	c	d3	e3	f3

Apply $A \rightarrow B$: $R_3(A=a) \rightarrow R_3.B = R_1.B = b$ Apply $B \rightarrow F$: $R_1(B=b) \rightarrow R_1.F = R_3.F = f$ Apply $C \rightarrow D$: $R_2(C=c) \rightarrow R_2.D = R_3.D = d$ Apply $D \rightarrow E$: $R_2(D=d) \rightarrow R_2.E = R_3.E = e$

Row R3 now has: (a, b, c, d, e, f) --- all distinguished symbols

The decomposition is LOSSLESS. ✓

[ANSWER BOX]

Lossless ✓ --- Chase produces a row with all distinguished symbols

Q3) a) ACID Properties and Transaction States

[9]

ACID properties ensure reliable processing of database transactions:

1. **Atomicity: Transaction executes completely or not at all (all-or-nothing).**
 - *Example:* Fund transfer --- debit and credit must both succeed or both fail
2. **Consistency: Transaction brings database from one valid state to another, preserving all integrity constraints.**
 - *Example:* Total balance before and after transfer must remain the same
3. **Isolation: Concurrent transactions appear to execute in isolation (serial order).**
 - *Example:* Two simultaneous transfers from same account produce same result as sequential
4. **Durability: Once committed, transaction's changes persist even after system failure.**
 - *Example:* After commit confirmation, power loss doesn't lose the update

Transaction State Diagram:

```
stateDiagram-v2
    [*] --> Active : Begin
    Active --> Partially_Committed : Final statement executed
    Active --> Failed : Abort during execution
    Partially_Committed --> Committed : All changes written to DB
    Partially_Committed --> Failed : Write/validation error
    Failed --> Aborted : Rollback completed
    Aborted --> [*]
    Committed --> [*]
```

Q3) b) Conflict Serializability

[9]

Step 1: Identify conflicting operations (same data item, at least one write):

- R1(A) and W2(A) --- **conflict** (T1 → T2? No, T2 happens after on different op... let me recheck)
- Let me reorder by time:

Time	T1	T2	T3
t1	R(A)		
t2		W(A)	
t3	R(B)		
t4			R(A)
t5		R(B)	
t6	W(A)		
t7		W(B)	
t8			W(A)

Conflicting pairs:

- W2(A) and W3(A): $T2 \rightarrow T3$ ($t2 < t8$)
- W2(A) and R1(A): $T2 \rightarrow T1$ ($t2 < t6$... wait)
- R1(B) and W2(B): $T1 \rightarrow T2$ ($t3 < t7$)
- R1(A) at $t1$ vs W2(A) at $t2$: $T1 \rightarrow T2$
- W1(A) at $t6$ vs R3(A) at $t4$: R3 happens before W1, so no conflict
- W1(A) at $t6$ vs W3(A) at $t8$: $T1 \rightarrow T3$

Precedence graph:

- $T1 \rightarrow T2$ (R1B $\rightarrow$ W2B)
- $T2 \rightarrow T3$ (W2A $\rightarrow$ W3A)
- $T1 \rightarrow T3$ (W1A $\rightarrow$ W3A... no, W3A is after... wait)

Let me rebuild carefully:

Conflicts on A:

- R1(A) [$t1$] and W2(A) [$t2$]: $T1 \rightarrow T2$
- W2(A) [$t2$] and W3(A) [$t8$]: $T2 \rightarrow T3$
- R1(A) [$t1$] and W3(A) [$t8$]: $T1 \rightarrow T3$... actually W3 happens after both R1 and W2

Conflicts on B:

- R1(B) [$t3$] and W2(B) [$t7$]: $T1 \rightarrow T2$
- R2(B) [$t5$] and W2(B) [$t7$]: $T2 \rightarrow T2$ (self --- ignore)

Precedence Graph:

$T1 \longrightarrow T2 \longrightarrow T3$

No cycles. The schedule is conflict serializable.

Equivalent serial schedule: $T1 \rightarrow T2 \rightarrow T3$

[ANSWER BOX]

Conflict Serializable ✓

Equivalent Serial Schedule: $\langle T1, T2, T3 \rangle$

Q4) a) Timestamp-Based Concurrency Control

[9]

Timestamp-based protocol uses timestamps to order transactions and ensure serializability without locking.

- **TS(T)** --- unique timestamp assigned to each transaction when it begins
- **R-timestamp(Q)** --- largest timestamp of any transaction that successfully read Q
- **W-timestamp(Q)** --- largest timestamp of any transaction that successfully wrote Q

Read Operation (Transaction T with timestamp TS):

- If $TS(T) < W\text{-timestamp}(Q)$, reject read and rollback T (data was written by a "later" transaction)
- Otherwise, execute read, set $R\text{-timestamp}(Q) = \max(R\text{-timestamp}(Q), TS(T))$

Write Operation (Transaction T with timestamp TS):

- If $TS(T) < R\text{-timestamp}(Q)$, reject write and rollback T (a younger transaction already read the old value)
- If $TS(T) < W\text{-timestamp}(Q)$, reject write (Thomas Write Rule --- can ignore outdated write)
- Otherwise, execute write, set $W\text{-timestamp}(Q) = TS(T)$

Example:

- Schedule: W1(A), R2(A) where $TS(T1)=10$, $TS(T2)=20$
- W1(A): $TS(T1)=10 \geq W-ts(A)=0 \rightarrow$ allowed. $W-ts(A)=10$
- R2(A): $TS(T2)=20 \geq W-ts(A)=10 \rightarrow$ allowed. $R-ts(A)=20 \checkmark$

Q4) b) Deadlock

[9]

Deadlock is a state where each transaction in a set is waiting for another transaction in the same set to release a lock.

Deadlock Prevention:

1. **Wait-Die:** Older transaction waits, younger dies (restarts with same timestamp)
 - If T1 (older) needs lock held by T2 (younger) $\rightarrow$ T1 waits
 - If T2 (younger) needs lock held by T1 (older) $\rightarrow$ T2 dies
2. **Wound-Wait:** Older wounds (preempts) younger, younger waits
 - If T1 (older) needs lock held by T2 (younger) $\rightarrow$ T1 wounds T2
 - If T2 (younger) needs lock held by T1 (older) $\rightarrow$ T2 waits

Deadlock Detection:

- Build **wait-for graph**: nodes = transactions, edges = $T_1 \rightarrow T_2$ if T_1 waits for T_2
- Periodically check for cycles; if cycle exists, **deadlock detected**
- Select a **victim** transaction (youngest/least cost) and roll it back

Deadlock Avoidance (Banker's Algorithm):

- Transaction declares maximum resource needs upfront
- System checks if granting a request leads to a **safe state** (where all transactions can finish)
- Only grant if state remains safe

Q5) a) CAP Theorem

[9]

CAP Theorem (Brewer's) states that a distributed data system can guarantee at most two of three properties:

1. **Consistency (C):** All nodes see the same data simultaneously. Any read returns the most recent write.
2. **Availability (A):** Every request receives a response (success or failure), regardless of node failures.
3. **Partition Tolerance (P):** System continues to operate despite network partitions (messages lost/delayed).

Trade-offs in NoSQL Databases:

System Type	CA	CP	AP
Behavior	Sacrifices partition tolerance	Sacrifices availability	Sacrifices consistency
Examples	Traditional RDBMS (single-node)	HBase, MongoDB	Cassandra, CouchDB
Use Case	Single-site transactional	Banking, financial	Social media, IoT

Key insight: In distributed systems, network partitions are inevitable, so P must be chosen. The real choice is CP vs AP.

Q5) b) RDBMS vs NoSQL and BASE Properties

[9]

Parameter	RDBMS	NoSQL
Data Model	Relational (tables/rows)	Key-value, document, graph, column
Schema	Fixed, rigid schema	Dynamic, schema-less
ACID Properties	Fully supported	BASE (Basically Available, Soft state, Eventual consistency)
Scalability	Vertical (scale-up)	Horizontal (scale-out)
Transactions	Full ACID, multi-row	Simple transactions, eventual consistency
Query Language	SQL (standardized)	Proprietary API per DB
Consistency	Strong consistency	Eventual consistency
Use Cases	Banking, ERP, finance	Big data, IoT, social media, real-time analytics

BASE Properties:

- **Basically Available:** System guarantees availability (CAP --- choose A)
- **Soft State:** State may change over time even without input (due to eventual consistency)
- **Eventual Consistency:** System becomes consistent over time; reads may not reflect latest writes immediately

Q6) a) NoSQL Data Models

[9]

1. Key-Value Store:

- Data stored as key-value pairs (like a hash map)
- *Examples:* Redis, DynamoDB, Riak
- *Use case:* Session management, caching

2. Document Store:

- Data stored as documents (JSON, BSON, XML)
- Documents have hierarchical structure
- *Examples:* MongoDB, CouchDB
- *Use case:* Content management, e-commerce catalogs
- *Document example:*

```
{
  "_id": "101",
  "name": "John Doe",
  "courses": ["DBMS", "TOC"],
  "address": { "city": "Pune", "zip": "411001" }
}
```

3. Column-Family Store:

- Data stored in column families (like Cassandra's column-oriented model)
- Each row can have different columns
- *Examples:* Cassandra, HBase
- *Use case:* Time-series data, IoT

4. Graph Database:

- Data stored as nodes (entities) and edges (relationships)
- *Examples:* Neo4j, OrientDB

- *Use case:* Social networks, recommendation engines

Q6) b) MongoDB CRUD Operations

[8]

1. Insert Operations:

```
// Insert one document
db.students.insertOne({ name: "Alice", age: 20, grade: "A" });

// Insert multiple documents
db.students.insertMany([
  { name: "Bob", age: 21, grade: "B" },
  { name: "Charlie", age: 19, grade: "A" },
]);
```

2. Query Operations:

```
// Find all documents
db.students.find();

// Find with filter
db.students.find({ grade: "A" });

// Find with operators
db.students.find({ age: { $gte: 20 } });

// Find one document
db.students.findOne({ name: "Alice" });
```

3. Update Operations:

```
// Update one document
db.students.updateOne({ name: "Alice" }, { $set: { grade: "A+" } });

// Update multiple documents
db.students.updateMany({ grade: "B" }, { $set: { status: "good" } });
```

4. Delete Operations:

```
// Delete one document
db.students.deleteOne({ name: "Charlie" });

// Delete multiple documents
db.students.deleteMany({ status: "graduated" });
```

Q7) a) Semi-Structured Data: JSON vs XML

[9]

Semi-structured data is data that does not conform to a rigid schema (like RDBMS tables) but has some organizational structure via tags or markers.

Features of Semi-Structured Data:

- Schema-on-read (structure is imposed when reading, not writing)
- Irregular or incomplete data allowed

- Self-describing (data contains its own schema/metadata)
- Nested/hierarchical representation

JSON vs XML:

Feature	JSON	XML
Full Form	JavaScript Object Notation	eXtensible Markup Language
Syntax	Lightweight, key-value pairs	Tag-based (opening/closing tags)
Data Types	String, Number, Boolean, Array, Object	All values are text (parsed by application)
Readability	More human-readable	Verbose
Parsing Speed	Fast	Slower
Namespace Support	No	Yes
Comments	Not supported natively	Supported
Use Case	Web APIs, REST services, MongoDB	Configuration files, SOAP, document storage

Q7) b) Object-Relational Database System and ORM

[9]

Object-Relational Database System (ORDBMS) extends the relational model with object-oriented features:


- User-defined types (UDTs)
- Inheritance
- Complex data types (nested tables, arrays)
- Methods/functions

Table Inheritance Example:

```
CREATE TYPE Person AS OBJECT (
  id NUMBER,
  name VARCHAR2(50),
  address VARCHAR2(100)
) NOT FINAL;

CREATE TYPE Student UNDER Person (
  roll_no NUMBER,
  grade VARCHAR2(2)
) FINAL;
```

Object-Relational Mapping (ORM) bridges object-oriented programming and relational databases by mapping classes to tables and objects to rows.



Benefits of ORM: Faster development, reduced boilerplate, database portability, caching.
Drawbacks: Performance overhead, complex queries harder to optimize.

Q8) a) Short Notes on Emerging Databases

[9]

1. Active Databases:

- Automatically respond to events using **ECA rules** (Event-Condition-Action)
- *Example:* Trigger in SQL --- `CREATE TRIGGER update_stock AFTER INSERT ON orders...`
- *Use:* Inventory management, audit logging

2. Deductive Databases:

- Use **logic programming** (Datalog) to derive new facts from existing ones
- Consist of facts (EDB --- extensional DB) and rules (IDB --- intensional DB)
- *Example:* `ancestor(X, Z) :- parent(X, Y), ancestor(Y, Z)`
- *Use:* Expert systems, genealogy

3. Spatial Databases:

- Store and query geometric/spatial data (points, lines, polygons)
- Support spatial indexing (R-trees) and spatial queries (nearby, inside, intersect)
- *Example:* PostGIS --- find all restaurants within 5 km
- *Use:* GIS, location-based services, navigation

4. Main Memory Databases:

- Store entire database in RAM for ultra-fast access
- Avoid disk I/O overhead; data persists via snapshots and logging
- *Examples:* Redis, SAP HANA, VoltDB
- *Use:* Real-time analytics, caching, session stores, gaming leaderboards

Q8) b) Nested Data Types --- JSON Document

[8]

JSON (JavaScript Object Notation) --- lightweight data-interchange format using key-value pairs.

XML (eXtensible Markup Language) --- markup language using custom tags for data representation.

JSON Document --- University System:

```
{
  "university": "Savitribai Phule Pune University",
  "location": "Pune, India",
  "departments": [
    {
      "name": "Computer Engineering",
      "code": "CO",
      "courses": [
        { "code": "310241", "name": "DBMS", "credits": 3 },
        { "code": "310242", "name": "TOC", "credits": 3 }
      ],
      "students": [
        {
          "id": "TE101",
          "name": "Alice",
          "semester": 5,
          "enrolled_courses": ["310241", "310242"]
        }
      ]
    },
    {
      "name": "Electronics Engineering",
      "code": "E&TC",
      "courses": [{ "code": "310261", "name": "DSP", "credits": 3 }]
    }
  ],
  "total_students": 1200,
  "established_year": 1949
}
```

[ANSWER BOX]

Key difference: JSON uses key-value pairs (lighter, faster to parse);
XML uses tag-based hierarchy (more verbose, supports namespaces and attributes)

EXAMINER COMMENTARY

Why this scores full marks:

- Each answer uses bold technical terms on first mention
- Q1(b) shows step-by-step closure computation with final boxed answer
- Q3(b) includes complete precedence graph analysis with timing trace
- Tables used wherever comparison was asked (Q5b, Q6a, Q7a)
- Mermaid diagram for transaction states and ASCII for precedence
- Practical MongoDB syntax in Q6(b)

Common Deductions:

- Not showing closure computation steps in normalization problems
- Missing arrow directions in precedence graphs
- Writing ACID definitions without real-world examples
- Confusing recoverable vs non-recoverable schedules

- Forgetting to check both BCNF and 3NF conditions

Time Budget:

- Q1 (18 min): 9 min theory + 9 min numerical
 - Q2 (18 min): 8 min theory + 9 min numerical
 - Q3 (18 min): 9 min ACID + 9 min serializability
 - Q4 (18 min): 9 min timestamp + 9 min deadlock
 - Q5 (18 min): 9 min CAP + 9 min comparison
 - Q6 (17 min): 9 min NoSQL types + 8 min MongoDB
 - Q7 (18 min): 9 min semi-structured + 9 min ORDBMS
 - Q8 (17 min): 9 min short notes + 8 min JSON
 - **Total: ~142 min** (within 150 min limit)
-
-
-